



Diffusion Generative AI for Computer Vision and Science

**—— Lecture 6: Machine Learning Frameworks and
Tools**

**Teaching Assistants: Kaihui Cheng
AI institute, Fudan University**

Contents

- **Introduce to Pytorch**
- **Project Guidelines**
 - **Project Expectations**
 - **Diffusion Demo**
 - **Picking a project idea**
 - **Final Report**

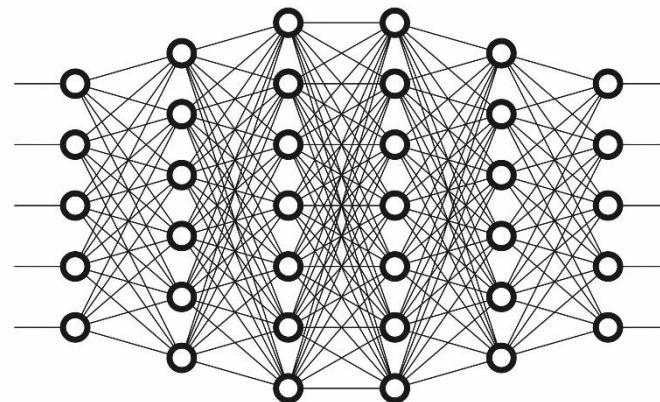
Pytorch

- Pipeline to train your model.

Training Data



Neural Network



Ground Truth

Loss
Function

Optimizer

Initializing a Tensor

Directly from data

Tensors can be created directly from data. The data type is automatically inferred.

```
data = [[1, 2],[3, 4]]
x_data = torch.tensor(data)
```

From a NumPy array

Tensors can be created from NumPy arrays (and vice versa - see [Bridge with NumPy](#)).

```
np_array = np.array(data)
x_np = torch.from_numpy(np_array)
```

From another tensor:

The new tensor retains the properties (shape, datatype) of the argument tensor, unless explicitly overridden.

```
x_ones = torch.ones_like(x_data) # retains the properties of x_data
print(f"Ones Tensor: \n {x_ones} \n")

x_rand = torch.rand_like(x_data, dtype=torch.float) # overrides the datatype
print(f"Random Tensor: \n {x_rand} \n")
```

Out:

Ones Tensor:
`tensor([[1, 1],
[1, 1]])`

Random Tensor:
`tensor([[0.8823, 0.9150],
[0.3829, 0.9593]])`

Attributes of a Tensor

Tensor attributes describe their shape, datatype, and the device on which they are stored.

```
tensor = torch.rand(3,4)

print(f"Shape of tensor: {tensor.shape}")
print(f"Datatype of tensor: {tensor.dtype}")
print(f"Device tensor is stored on: {tensor.device}")
```

Out:

```
Shape of tensor: torch.Size([3, 4])
Datatype of tensor: torch.float32
Device tensor is stored on: cpu
```

Operations on Tensors

By default, tensors are created on the CPU. We need to explicitly move tensors to the GPU using `.to` method (after checking for GPU availability). Keep in mind that copying large tensors across devices can be expensive in terms of time and memory!

```
# We move our tensor to the GPU if available
if torch.cuda.is_available():
    tensor = tensor.to("cuda")
```

Operations on Tensors

Standard numpy-like indexing and slicing:

```
tensor = torch.ones(4, 4)
print(f"First row: {tensor[0]}")
print(f"First column: {tensor[:, 0]}")
print(f"Last column: {tensor[..., -1]}")
tensor[:, 1] = 0
print(tensor)
```

Out:

```
First row: tensor([1., 1., 1., 1.])
First column: tensor([1., 1., 1., 1.])
Last column: tensor([1., 1., 1., 1.])
tensor([[1., 0., 1., 1.],
        [1., 0., 1., 1.],
        [1., 0., 1., 1.],
        [1., 0., 1., 1.]])
```

Arithmetic operations

```
# This computes the matrix multiplication between two tensors. y1, y2, y3 will
# have the same value
# ``tensor.T`` returns the transpose of a tensor
```

```
y1 = tensor @ tensor.T
y2 = tensor.matmul(tensor.T)
```

```
y3 = torch.rand_like(y1)
torch.matmul(tensor, tensor.T, out=y3)
```

```
# This computes the element-wise product. z1, z2, z3 will have the same value
```

```
z1 = tensor * tensor
z2 = tensor.mul(tensor)
```

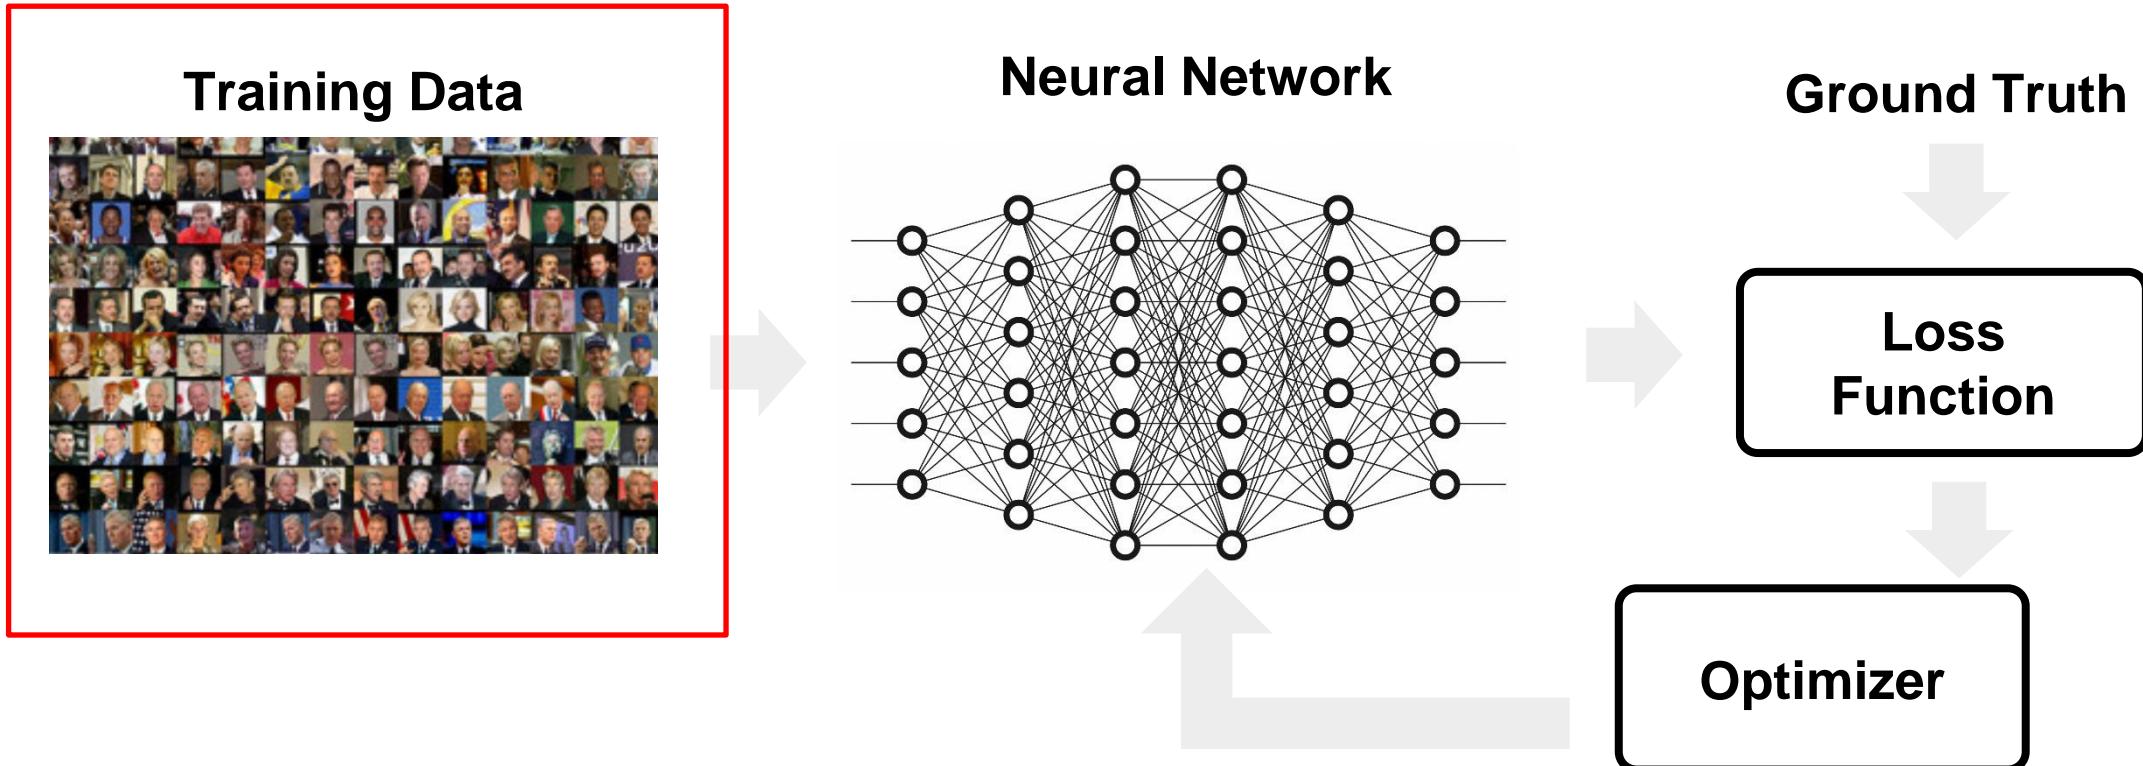
```
z3 = torch.rand_like(tensor)
torch.mul(tensor, tensor, out=z3)
```

Out:

```
tensor([[1., 0., 1., 1.],
        [1., 0., 1., 1.],
        [1., 0., 1., 1.],
        [1., 0., 1., 1.]])
```

Pytorch

- Pipeline to train your model.



Loading a Dataset

```
import torch
from torch.utils.data import Dataset
from torchvision import datasets
from torchvision.transforms import ToTensor
import matplotlib.pyplot as plt

training_data = datasets.FashionMNIST(
    root="data",
    train=True,
    download=True,
    transform=ToTensor()
)

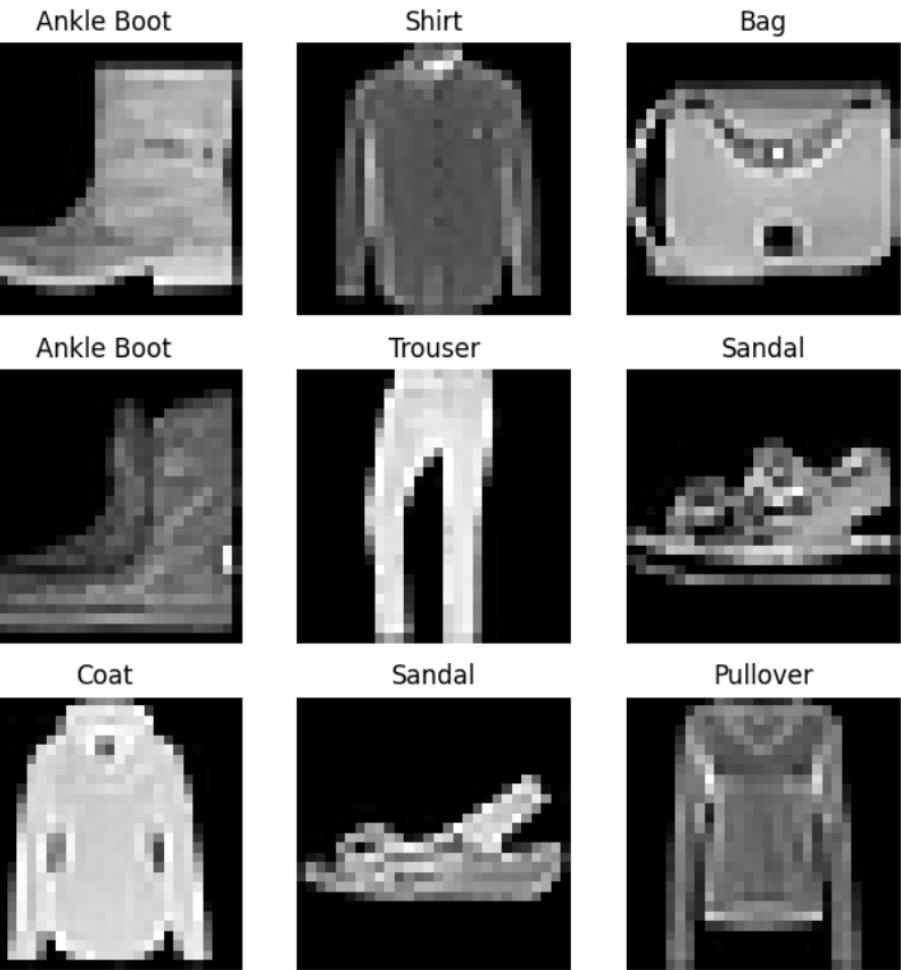
test_data = datasets.FashionMNIST(
    root="data",
    train=False,
    download=True,
    transform=ToTensor()
)
```

Iterating and Visualizing the Dataset

```

labels_map = {
    0: "T-Shirt",
    1: "Trouser",
    2: "Pullover",
    3: "Dress",
    4: "Coat",
    5: "Sandal",
    6: "Shirt",
    7: "Sneaker",
    8: "Bag",
    9: "Ankle Boot",
}
figure = plt.figure(figsize=(8, 8))
cols, rows = 3, 3
for i in range(1, cols * rows + 1):
    sample_idx = torch.randint(len(training_data), size=(1,)).item()
    img, label = training_data[sample_idx]
    figure.add_subplot(rows, cols, i)
    plt.title(labels_map[label])
    plt.axis("off")
    plt.imshow(img.squeeze(), cmap="gray")
plt.show()

```



Creating a Custom Dataset for your files

```
import os
import pandas as pd
from torchvision.io import read_image

class CustomImageDataset(Dataset):
    def __init__(self, annotations_file, img_dir, transform=None, target_transform=None):
        self.img_labels = pd.read_csv(annotations_file)
        self.img_dir = img_dir
        self.transform = transform
        self.target_transform = target_transform

    def __len__(self):
        return len(self.img_labels)

    def __getitem__(self, idx):
        img_path = os.path.join(self.img_dir, self.img_labels.iloc[idx, 0])
        image = read_image(img_path)
        label = self.img_labels.iloc[idx, 1]
        if self.transform:
            image = self.transform(image)
        if self.target_transform:
            label = self.target_transform(label)
        return image, label
```

Creating a Custom Dataset for your files

```
import os
import pandas as pd
from torchvision.io import read_image

class CustomImageDataset(Dataset):
    def __init__(self, annotations_file, img_dir, transform=None, target_transform=None):
        self.img_labels = pd.read_csv(annotations_file)
        self.img_dir = img_dir
        self.transform = transform
        self.target_transform = target_transform

    def __len__(self):
        return len(self.img_labels)

    def __getitem__(self, idx):
        img_path = os.path.join(self.img_dir, self.img_labels.iloc[idx, 0])
        image = read_image(img_path)
        label = self.img_labels.iloc[idx, 1]
        if self.transform:
            image = self.transform(image)
        if self.target_transform:
            label = self.target_transform(label)
        return image, label
```

Creating a Custom Dataset for your files

```
import os
import pandas as pd
from torchvision.io import read_image

class CustomImageDataset(Dataset):
    def __init__(self, annotations_file, img_dir, transform=None, target_transform=None):
        self.img_labels = pd.read_csv(annotations_file)
        self.img_dir = img_dir
        self.transform = transform
        self.target_transform = target_transform

    def __len__(self):
        return len(self.img_labels)

    def __getitem__(self, idx):
        img_path = os.path.join(self.img_dir, self.img_labels.iloc[idx, 0])
        image = read_image(img_path)
        label = self.img_labels.iloc[idx, 1]
        if self.transform:
            image = self.transform(image)
        if self.target_transform:
            label = self.target_transform(label)
        return image, label
```

Creating a Custom Dataset for your files

```
import os
import pandas as pd
from torchvision.io import read_image

class CustomImageDataset(Dataset):
    def __init__(self, annotations_file, img_dir, transform=None, target_transform=None):
        self.img_labels = pd.read_csv(annotations_file)
        self.img_dir = img_dir
        self.transform = transform
        self.target_transform = target_transform

    def __len__(self):
        return len(self.img_labels)

    def __getitem__(self, idx):
        img_path = os.path.join(self.img_dir, self.img_labels.iloc[idx, 0])
        image = read_image(img_path)
        label = self.img_labels.iloc[idx, 1]
        if self.transform:
            image = self.transform(image)
        if self.target_transform:
            label = self.target_transform(label)
        return image, label
```

Preparing your data for training with DataLoaders

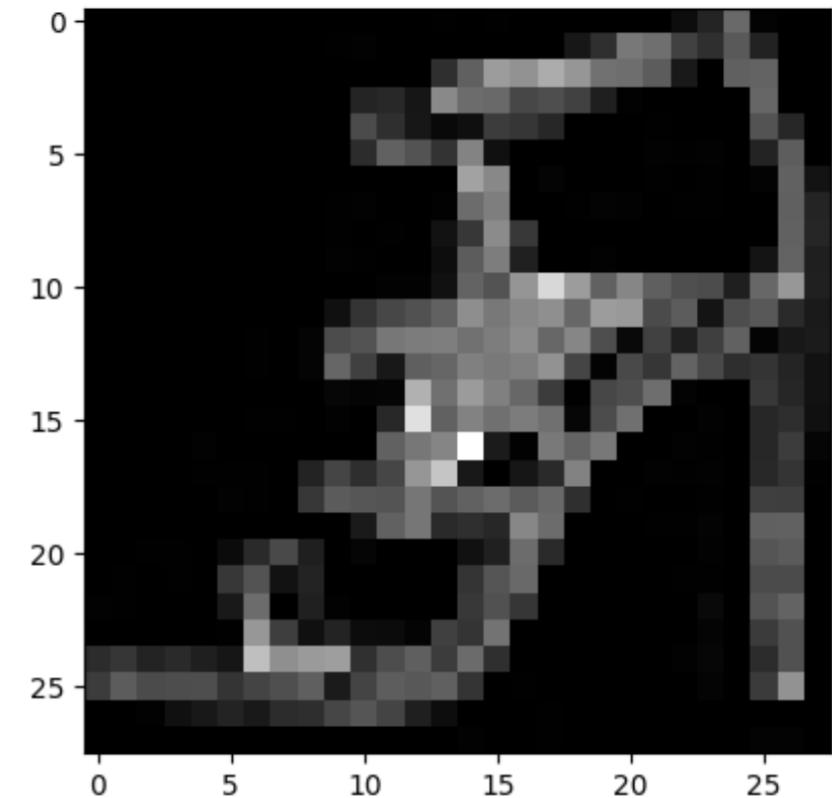
The `Dataset` retrieves our dataset's features and labels one sample at a time. While training a model, we typically want to pass samples in “minibatches”, reshuffle the data at every epoch to reduce model overfitting, and use Python’s `multiprocessing` to speed up data retrieval.

`DataLoader` is an iterable that abstracts this complexity for us in an easy API.

```
from torch.utils.data import DataLoader

train_dataloader = DataLoader(training_data, batch_size=64, shuffle=True)
test_dataloader = DataLoader(test_data, batch_size=64, shuffle=True)
```

```
# Display image and label.
train_features, train_labels = next(iter(train_dataloader))
print(f"Feature batch shape: {train_features.size()}")
print(f"Labels batch shape: {train_labels.size()}")
img = train_features[0].squeeze()
label = train_labels[0]
plt.imshow(img, cmap="gray")
plt.show()
print(f"Label: {label}")
```



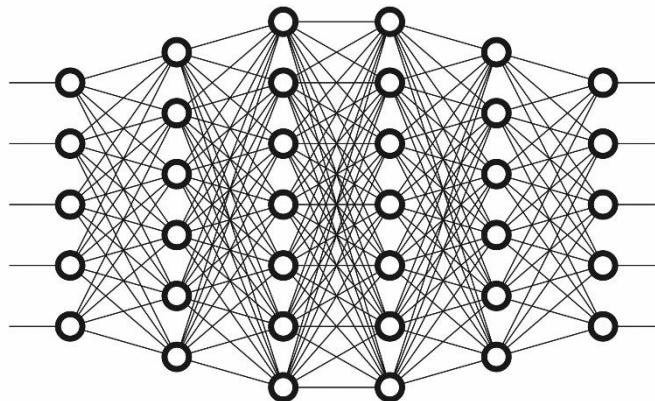
Pytorch

- Pipeline to train your model.

Training Data



Neural Network



Ground Truth

Loss
Function

Optimizer

Build a Neural Network

```
import os
import torch
from torch import nn
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
```

We want to be able to train our model on a hardware accelerator like the GPU or MPS, if available. Let's check to see if `torch.cuda` or `torch.backends.mps` are available, otherwise we use the CPU.

```
device = (
    "cuda"
    if torch.cuda.is_available()
    else "mps"
    if torch.backends.mps.is_available()
    else "cpu"
)
print(f"Using {device} device")
```

Build a Neural Network

We define our neural network by subclassing `nn.Module`, and initialize the neural network layers in `__init__`. Every `nn.Module` subclass implements the operations on input data in the `forward` method.

```
class NeuralNetwork(nn.Module):
    def __init__(self):
        super().__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28*28, 512),
            nn.ReLU(),
            nn.Linear(512, 512),
            nn.ReLU(),
            nn.Linear(512, 10),
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits
```

Build a Neural Network

We create an instance of `NeuralNetwork`, and move it to the `device`, and print its structure.

```
model = NeuralNetwork().to(device)
print(model)
```

Out:

```
NeuralNetwork(
  (flatten): Flatten(start_dim=1, end_dim=-1)
  (linear_relu_stack): Sequential(
    (0): Linear(in_features=784, out_features=512, bias=True)
    (1): ReLU()
    (2): Linear(in_features=512, out_features=512, bias=True)
    (3): ReLU()
    (4): Linear(in_features=512, out_features=10, bias=True)
  )
)
```

To use the model, we pass it the input data. This executes the model's `forward`, along with some **background operations**. Do not call `model.forward()` directly!

Calling the model on the input returns a 2-dimensional tensor with dim=0 corresponding to each output of 10 raw predicted values for each class, and dim=1 corresponding to the individual values of each output. We get the prediction probabilities by passing it through an instance of the `nn.Softmax` module.

```
X = torch.rand(1, 28, 28, device=device)
logits = model(X)
pred_probab = nn.Softmax(dim=1)(logits)
y_pred = pred_probab.argmax(1)
print(f"Predicted class: {y_pred}")
```

Out:

```
Predicted class: tensor([7], device='cuda:0')
```

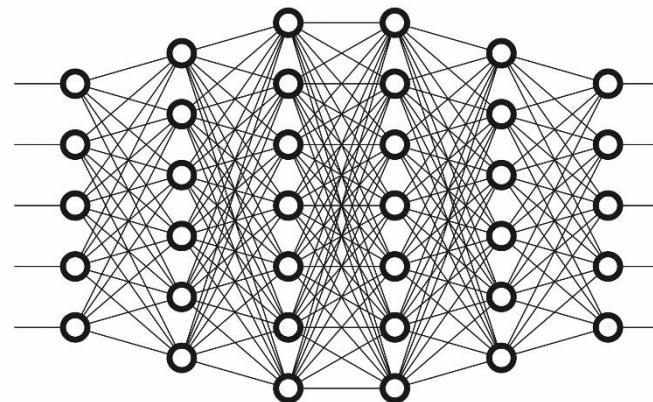
Pytorch

- Pipeline to train your model.

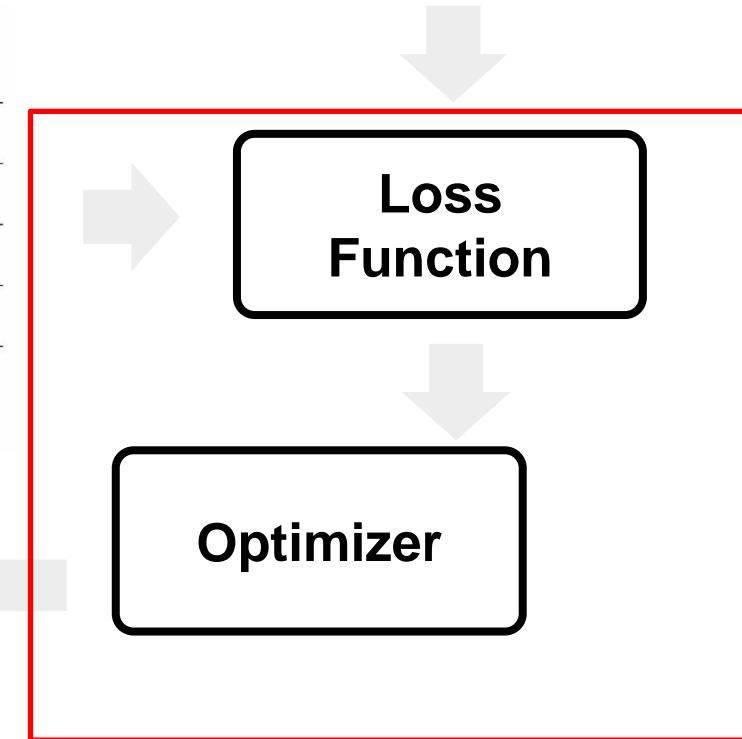
Training Data



Neural Network



Ground Truth



Automatic Differentiation

```
import torch

x = torch.ones(5) # input tensor
y = torch.zeros(3) # expected output
w = torch.randn(5, 3, requires_grad=True)
b = torch.randn(3, requires_grad=True)
z = torch.matmul(x, w)+b
loss = torch.nn.functional.binary_cross_entropy_with_logits(z, y)
```

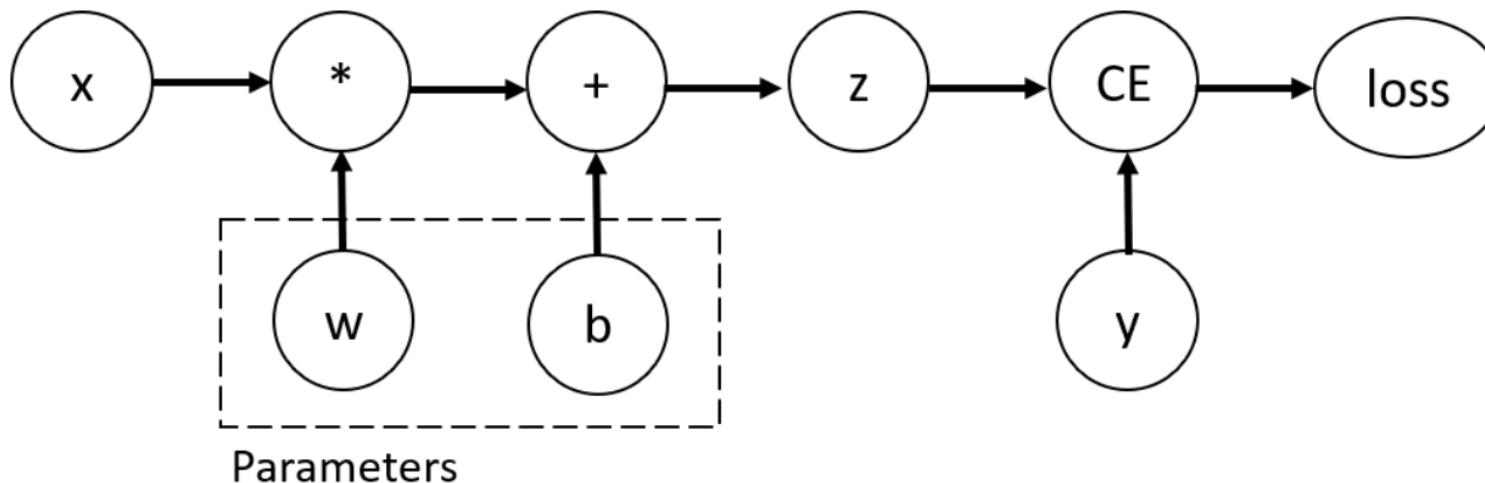
```
print(f"Gradient function for z = {z.grad_fn}")
print(f"Gradient function for loss = {loss.grad_fn}")
```

Out:

Gradient function for z = <AddBackward0 object at 0x7f55696a38b0>
 Gradient function for loss = <BinaryCrossEntropyWithLogitsBackward0 object at 0x7f55696a0e20>

Tensors, Functions and Computational graph

This code defines the following **computational graph**:

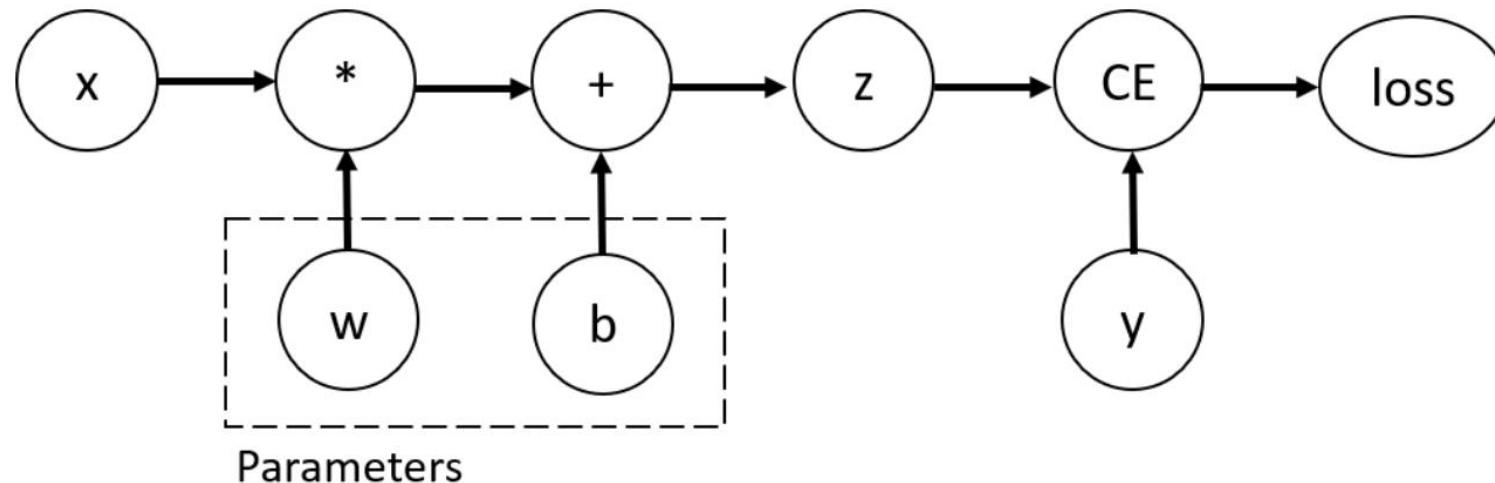


```
loss.backward()
print(w.grad)
print(b.grad)
```

Out:

tensor([[0.3313, 0.0626, 0.2530],
 [0.3313, 0.0626, 0.2530],
 [0.3313, 0.0626, 0.2530],
 [0.3313, 0.0626, 0.2530],
 [0.3313, 0.0626, 0.2530]])
tensor([0.3313, 0.0626, 0.2530])

Automatic Differentiation



```

loss.backward()
print(w.grad)
print(b.grad)
  
```

```

print(f"Gradient function for z = {z.grad_fn}")
print(f"Gradient function for loss = {loss.grad_fn}")
  
```

Out:

```

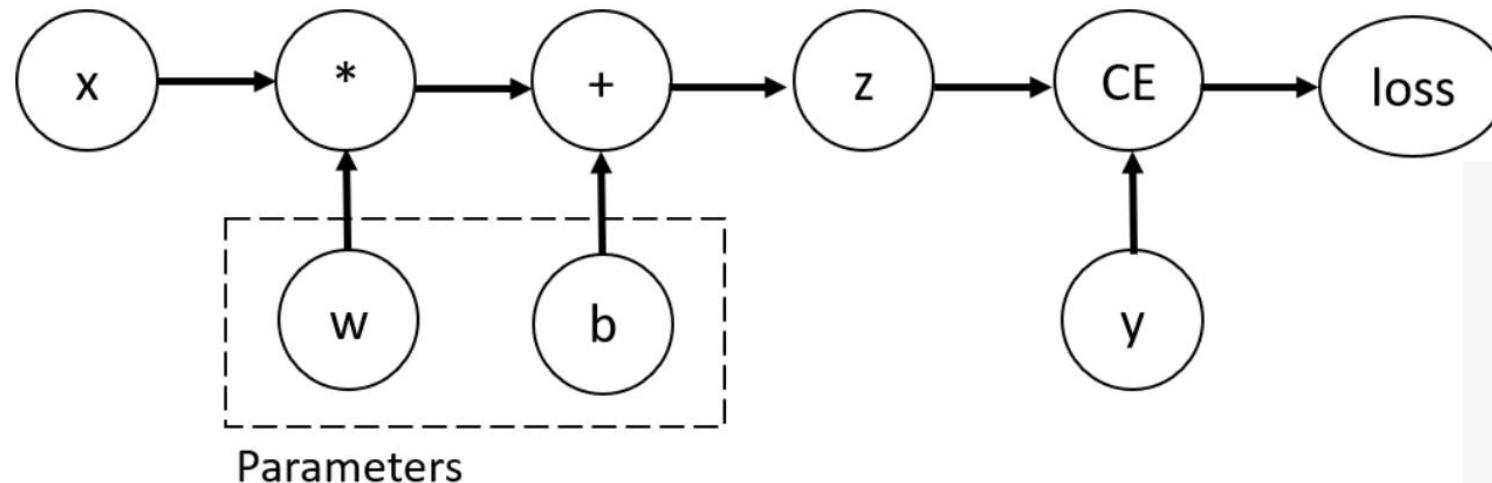
Gradient function for z = <AddBackward0 object at 0x7f55696a38b0>
Gradient function for loss = <BinaryCrossEntropyWithLogitsBackward0
object at 0x7f55696a0e20>
  
```

Out:

```

tensor([[0.3313, 0.0626, 0.2530],
       [0.3313, 0.0626, 0.2530],
       [0.3313, 0.0626, 0.2530],
       [0.3313, 0.0626, 0.2530],
       [0.3313, 0.0626, 0.2530]])
tensor([0.3313, 0.0626, 0.2530])
  
```

Automatic Differentiation



```

z = torch.matmul(x, w)+b
z_det = z.detach()
print(z_det.requires_grad)
  
```

Out:

False

```

z = torch.matmul(x, w)+b
print(z.requires_grad)
  
```

```

with torch.no_grad():
    z = torch.matmul(x, w)+b
print(z.requires_grad)
  
```

Out:

True
False

Loss Function & Optimizer

```

class NeuralNetwork(nn.Module):
    def __init__(self):
        super().__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28*28, 512),
            nn.ReLU(),
            nn.Linear(512, 512),
            nn.ReLU(),
            nn.Linear(512, 10),
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits
  
```

```

# Initialize the loss function
loss_fn = nn.CrossEntropyLoss()
  
```



```

optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)
  
```

Inside the training loop, optimization happens in three steps:

- Call `optimizer.zero_grad()` to reset the gradients of model parameters. Gradients by default add up; to prevent double-counting, we explicitly zero them at each iteration.
- Backpropagate the prediction loss with a call to `loss.backward()`. PyTorch deposits the gradients of the loss w.r.t. each parameter.
- Once we have our gradients, we call `optimizer.step()` to adjust the parameters by the gradients collected in the backward pass.

Train Loop

```
def train_loop(dataloader, model, loss_fn, optimizer):
    size = len(dataloader.dataset)
    # Set the model to training mode - important for batch normalization and dropout layers
    # Unnecessary in this situation but added for best practices
    model.train()
    for batch, (X, y) in enumerate(dataloader):
        # Compute prediction and loss
        pred = model(X)
        loss = loss_fn(pred, y)

        # Backpropagation
        loss.backward()
        optimizer.step()
        optimizer.zero_grad()

        if batch % 100 == 0:
            loss, current = loss.item(), batch * batch_size + len(X)
            print(f"loss: {loss:.7f} [{current:.5d}/{size:.5d}]")
```

Test Loop

```
def test_loop(dataloader, model, loss_fn):
    # Set the model to evaluation mode - important for batch normalization and dropout layers
    # Unnecessary in this situation but added for best practices
    model.eval()
    size = len(dataloader.dataset)
    num_batches = len(dataloader)
    test_loss, correct = 0, 0

    # Evaluating the model with torch.no_grad() ensures that no gradients are computed during
    # test mode
    # also serves to reduce unnecessary gradient computations and memory usage for tensors with
    # requires_grad=True
    with torch.no_grad():
        for X, y in dataloader:
            pred = model(X)
            test_loss += loss_fn(pred, y).item()
            correct += (pred.argmax(1) == y).type(torch.float).sum().item()

    test_loss /= num_batches
    correct /= size
    print(f"Test Error: \n Accuracy: {(100*correct):>0.1f}%, Avg loss: {test_loss:>8f} \n")
```

Full Implementation

```
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)

epochs = 10
for t in range(epochs):
    print(f"Epoch {t+1}\n-----")
    train_loop(train_dataloader, model, loss_fn, optimizer)
    test_loop(test_dataloader, model, loss_fn)
print("Done!")
```

Out:

```
Epoch 1
-----
loss: 2.298730 [ 64/60000]
loss: 2.289123 [ 6464/60000]
loss: 2.273286 [12864/60000]
loss: 2.269406 [19264/60000]
loss: 2.249603 [25664/60000]
```

Save & Load the Model Weights

PyTorch models store the learned parameters in an internal state dictionary, called `state_dict`. These can be persisted via the `torch.save` method:

```
model = models.vgg16(weights='IMAGENET1K_V1')
torch.save(model.state_dict(), 'model_weights.pth')
```

In the code below, we set `weights_only=True` to limit the functions executed during unpickling to only those necessary for loading weights. Using `weights_only=True` is considered a best practice when loading weights.

```
model = models.vgg16() # we do not specify ``weights'', i.e. create untrained model
model.load_state_dict(torch.load('model_weights.pth', weights_only=True))
model.eval()
```

Save & Load the Model with Shapes

When loading model weights, we needed to instantiate the model class first, because the class defines the structure of a network. We might want to save the structure of this class together with the model, in which case we can pass `model` (and not `model.state_dict()`) to the saving function:

```
torch.save(model, 'model.pth')
```

We can then load the model as demonstrated below.

As described in [Saving and loading torch.nn.Modules](#), saving `state_dict`'s is considered the best practice. However, below we use `weights_only=False` because this involves loading the model, which is a legacy use case for `torch.save`.

```
model = torch.load('model.pth', weights_only=False),
```

Contents

- Introduce to Pytorch
- Project Guidelines
 - **Project Expectations**
 - Diffusion Demo
 - Picking a project idea
 - Final Report

Course Description

- **Objective**

- Master the principles and theories of diffusion model-based generative AI.
- Build programming skills and project development experience.
- Apply diffusion-based generative AI to real-world problems.

- **Coursework**

- Attendance
- Participation
- Diffusion Demo
- Course Project
- **Students are encouraged to develop innovative projects that combine generative models with their academic disciplines.**

Project Expectations

- **Explore concepts on a topic of your choice with diffusion**
 - We will build a project page, see it in the later notifications.
- **Completed in groups, no more than 3**
 - The number of teams can be adjusted according to the difficulty of the task, but the project expectation are higher with more people
- **Three tracks:**
 - **Survey:** you can explore existing literature and methodologies in the diffusion algorithm space, identifying gaps and opportunities for innovation.
 - **Applications:** you can apply diffusion algorithm to a specific interest (e.g. vision, biology, physics, geosciences), it is great to see solve problems on your particular domains.
 - **Models:** you can build a new model or algorithm, test and verify your design. This might be challenging, but could come up with a publishable work or works !!!

Final Reports

- The final reports (6-8 pages) need to include:
 - **Title:** The main theme of your work (method name/ domain/ tasks)
 - **Author(s):** List of all contributors
 - **Abstract:** Brief summary of your work (problem, method, key findings)
 - **Introduction:** Outline the objectives of your work, motivation and the problem you seek to solve.
 - **Related Work:** Discuss prior similar or related work, and what is the difference of your work.
 - **Methods:** Provide a detailed description of your methods

Final Reports

- The final reports (6-8 pages) need to include:
 - **Implementation:** Explain what datasets and tools/devices you used
 - **Experiments:** Outline the experimental design, and present your findings clearly with figures or tables.
 - **Conclusion:** Summarize the main findings of method.
 - **Supplementary Materials (optional)** : Include any additional data, code, or resources that support your method
 - **Poster:** Include title, authors, introduction, methods, results and conclusion in a visually appealing poster.

Final Reports

- The final reports (6-8 pages) need to include:
 - Title
 - Author(s)
 - Abstract
 - Introduction
 - Related Work
 - Methods
 - Implementation
 - Experiments
 - Conclusion
 - Supplementary Materials (optional)
 - Poster

Examples of Posters

¹ Google Research



² UNIVERSITY OF WASHINGTON

³

Sean Fanello¹

Motivation

- > Harsh lighting in photography can result in strong specularities and shadows. Professional photographers address the problem by putting a diffuser or scrim between the subject and the primary light source.



Uncontrolled Light



Diffuser / Scrim

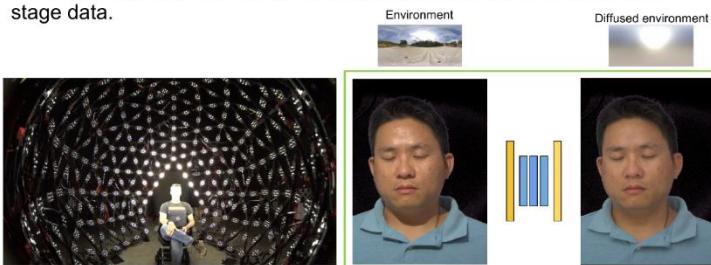


Diffused Light

- > We propose to provide the same light diffusion computationally, using just the photo with uncontrolled light.

Overview

We treat the problem as image-to-image translation supervised by light stage data.



Controllable Light Diffusion for Portraits

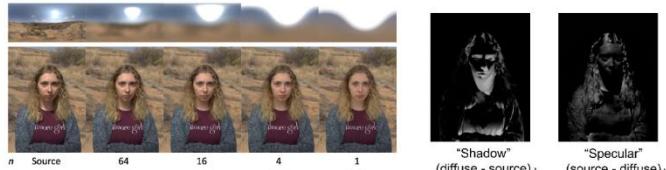
David Futschik^{1,2} Kelvin Ritland¹ James Vecore¹

Sergio Orts-Escalano¹ Brian Curless^{1,3} Daniel Sýkora^{1,2} Rohit Pandey¹

JUNE 18-22, 2023
CVPR VANCOUVER, CANADA

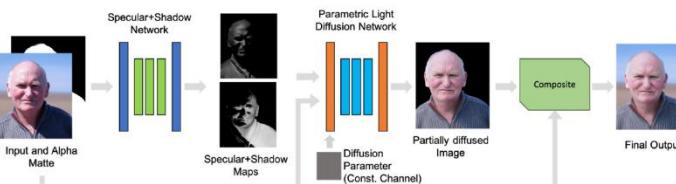
Method and Data

- > Use light stage to capture dataset of one light at a time images for relighting.
- > Relight subjects using high contrast HDR environment maps.
- > Use specular convolution with Phong exponent n to soften environment map lighting.
- > Represent the amount of diffusion by the Gini coefficient G to equalize the perceptual effect and normalize across different lighting environments.
$$G = \frac{\sum_{i=0}^m \sum_{j=0}^m |x_i - x_j|}{2m\bar{x}} \quad x_i - \text{env. pixel values}$$
- > Provide "shadow" and "specular" maps, where input image is darker or lighter than the fully diffused ($n=1$) image, resp., as additional supervision.

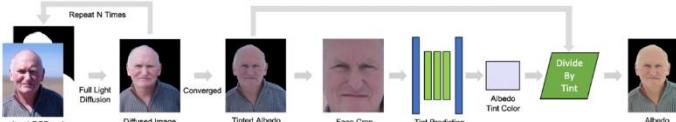


- > Generate dataset of 12 million examples: 61 identities, 270 HDR environment maps randomly rotated, random diffusion levels, augmentations including external shadows.

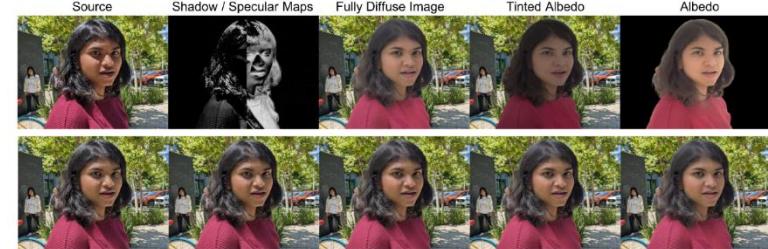
Light Diffusion Network



Extension to Albedo Prediction



Results



[1] Pandey et al.: Total Relighting: Learning to Relight Portraits for Background Replacement, SIGGRAPH 2021.
[2] Weir et al.: Deep Portrait Delighting, ECCV 2022.

[3] Yeh et al.: Learning to Relight Portrait Images via a Virtual Light Stage and Synthetic-to-Real Adaptation, SIGGRAPH Asia 2022.

Examples of Posters



Background

Imbalanced datasets are common in the medical imaging context where normal data outnumbers disease data. Deep learning models trained on imbalanced data tend to be biased towards majority class, leading to more false negatives, which is particularly problematic since this means medical conditions may be left undetected. Classic approaches include:

- Oversampling
- Undersampling
- Class weights
- SMOTE methods
- Data augmentation

To avoid overfitting that comes with classic oversampling methods we propose using a GAN to generate synthetic examples from random noise, adding to the diversity of minority class.

Experiments

	Precision	Recall/Sensitivity	Specificity	F1
Baseline	0.730	0.842	0.850	0.782
Class Weights	0.859	0.855	0.915	0.857
Under-sampling	0.709	0.865	0.843	0.779
Over-sampling	0.863	0.845	0.917	0.854
DCGAN	0.821	0.910	0.898	0.863
cDCGAN $\alpha = 0.50$	0.893	0.857	0.934	0.874
cDCGAN $\alpha = 0.25$	0.710	0.706	0.949	0.803
cDCGAN $\alpha = 0.75$	0.821	0.901	0.898	0.859

Problem

Our goal is to rebalance image datasets while reducing bias and improving overall model performance.

Our GAN uses random noise to generate new synthetic pneumonia chest X-rays that are added to the dataset, and we assess the different methods by training a ResNet-18 model that takes the chest X-rays as input and outputs a binary classification of healthy or pneumonia.

Due to our imbalanced dataset, we evaluate performance based not on accuracy, but precision, recall/sensitivity, specificity, and F1 score.

Visualizations

The visualizations show four pairs of chest X-ray images and their corresponding Grad-CAM heatmaps. The first pair is labeled "DCGAN Pneumonia" and "cDCGAN Healthy vs Pneumonia". The second pair is labeled "Baseline Grad-CAM". The third pair is labeled "DCGAN Grad-CAM". The fourth pair is labeled "cDCGAN Grad-CAM, $\alpha = 0.5$ ". Each heatmap shows the areas of the X-ray that were most influential for the model's classification decision.

Dataset

Dataset was taken from Kermany et al. consisting of pediatric chest X-rays from Guangzhou. We supplemented with healthy chest X-rays from public NIH dataset ChestX-Ray8. We randomly removed pneumonia images to create imbalanced situation: 4385 normal, 1460 pneumonia. The test set is more balanced to be an unbiased evaluation of model performance.

Analysis

We see that:

- DCGAN and the cDCGAN variants were all able to outperform the baseline in almost all metrics.
- In particular, both our DCGAN and cDCGAN ($\alpha = 0.5$) approaches outperform class weights, oversampling, and undersampling (the most commonly used methods) in terms of F1 score.
- DCGAN and cDCGAN ($\alpha = 0.25$) have the highest sensitivity and specificity, respectively. This uncovers the possibility of using these models in conjunction to generate a better dataset.
- We find that α can be used to tune sensitivity and specificity.
- Our Grad-CAMs uncover that the DCGAN ResNet is not learning relevant features inside the chest cavity, while the baseline and cDCGAN do a bit better in this respect.

Medical Image Dataset Rebalancing via Conditional Deep Convolutional Generative Adversarial Networks (cDCGANs)

Adam Sun and Kevin Li
 {adsun, kevinli8}@stanford.edu



Methods

We use a Generative Adversarial Network, which is used to learn a mapping from random noise to the data space of a distribution. A generator network, which tries to generate a realistic data point, plays a minimax game against a discriminator network which seeks to discriminate between real and fake data. The objective function from the original paper:

$$\min_{G} \max_{D} V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{data}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

We use a deep convolutional GAN (DCGAN), which has strided convolution, batchnorm, and ReLU to improve GAN stability and image comprehensibility. We train it on the minority class of pneumonia images and then use the generator to generate realistic examples from the minority class distribution and oversample the minority class to rebalance an imbalanced dataset.

Additionally, we explore conditional DCGAN as a way to utilize the entire dataset. By conditioning on labels, we can then drive the generated images towards the pneumonia class. We also introduce a hyperparameter α that represents the probability of a randomly generated healthy label.

Additionally, we compare ResNet-18 performance on the rebalanced dataset with its performance on datasets rebalanced by more classic methods.

CGAN architecture

```

    graph TD
      RN[Random Noise] --> G[Generator]
      EL[Encoded Labels] --> G
      G --> RDL[Real Data and Labels]
      RDL --> D[Discriminator]
      D --> RF[Real/Fake]
  
```

Conclusions and Future Work

In this work, we explored the usage of Deep Convolutional Generative Adversarial Networks (DCGANs) and conditional Deep Convolutional Generative Adversarial Networks (cDCGANs) for oversampling of the minority class in imbalanced medical image datasets.

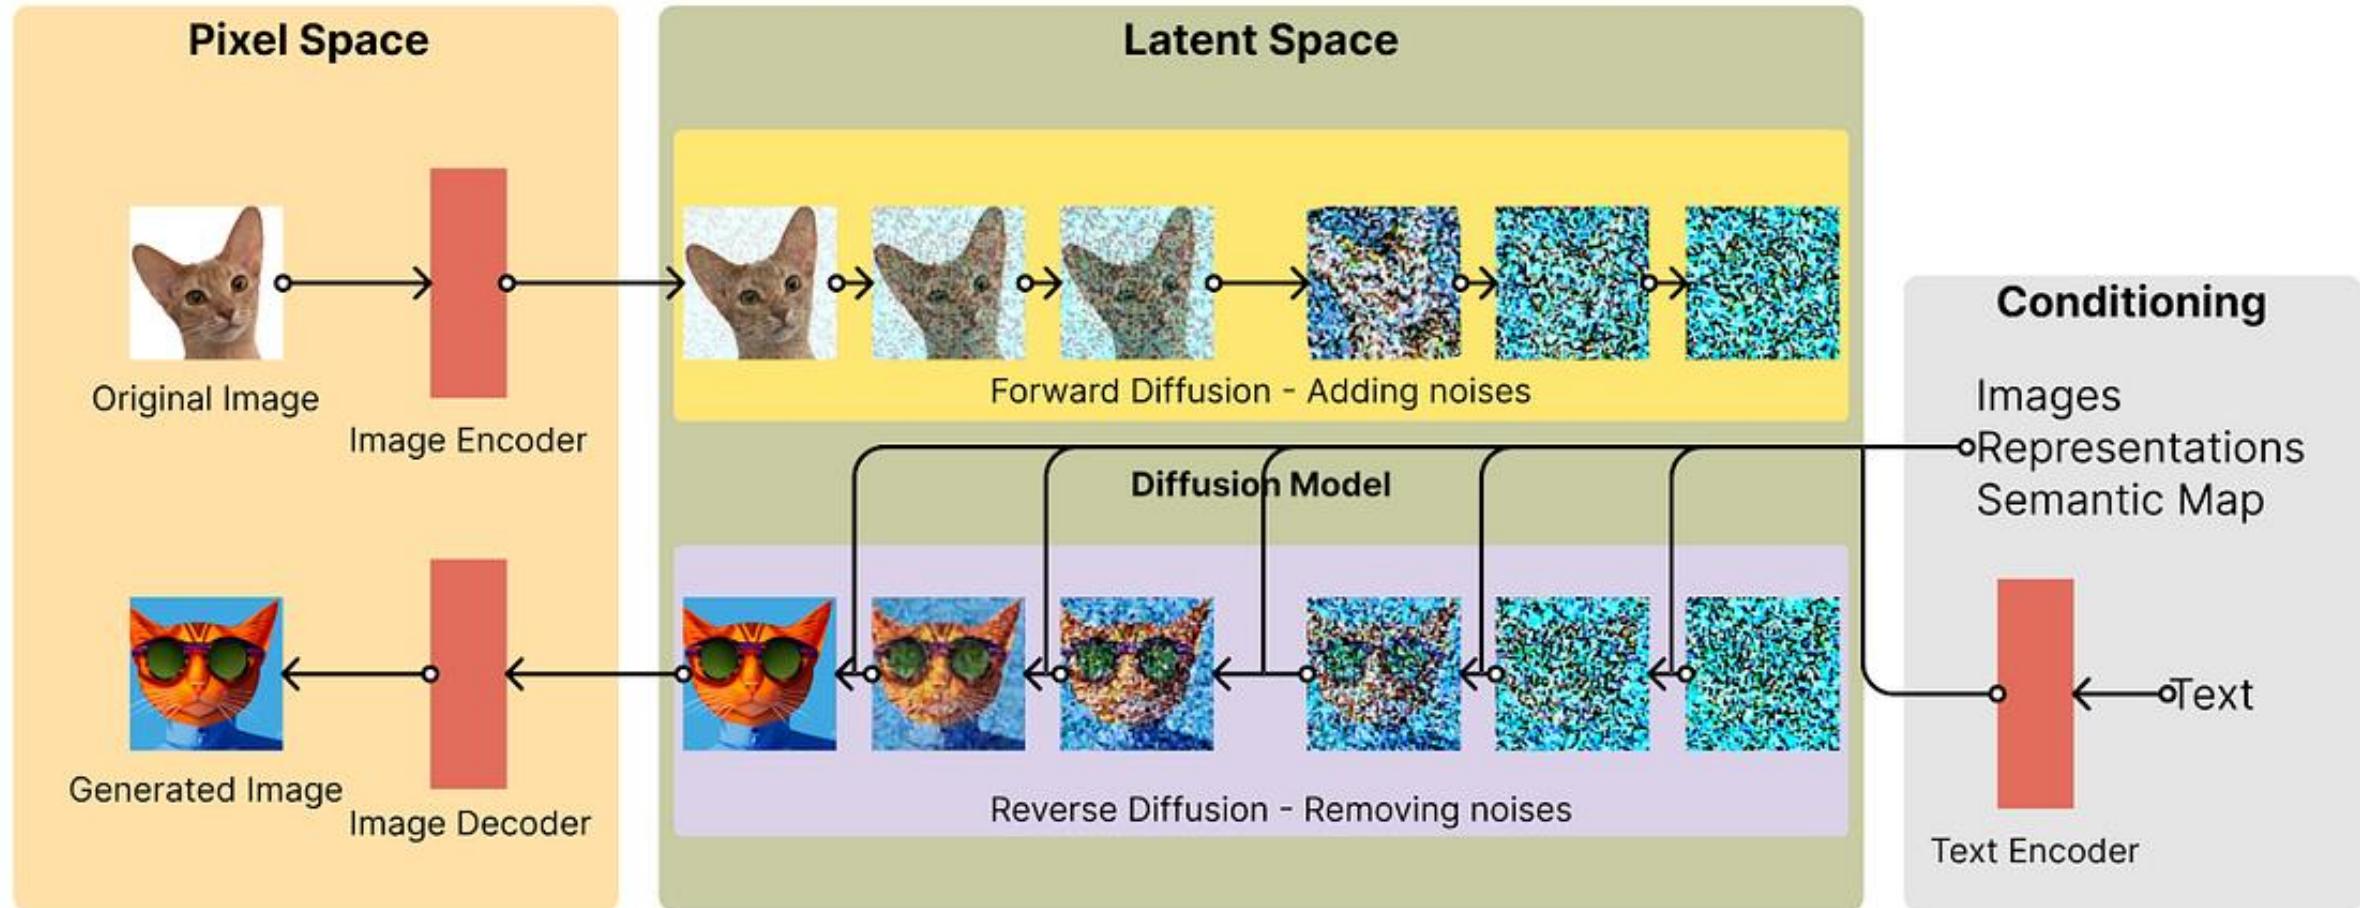
Based on the results, we conclude that GAN-based approaches are a valid way to oversample an imbalanced training dataset, despite the Grad-CAMs indicating that the model may be learning irrelevant features.

In the future, we would like to tune the GAN models more carefully, and train models multiple times to compensate for the stochastic nature of model training, as well as curate a more diverse imbalanced dataset. We may offer our GAN-generated images to an actual experienced radiologist for review for usefulness, relevance, and realism, resulting in a better annotated dataset.

Contents

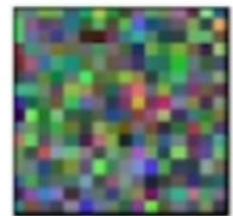
- Introduce to Pytorch
- Project Guidelines
 - Project Expectations
 - Diffusion Demo
 - Picking a project idea
 - Final Report

Diffusion



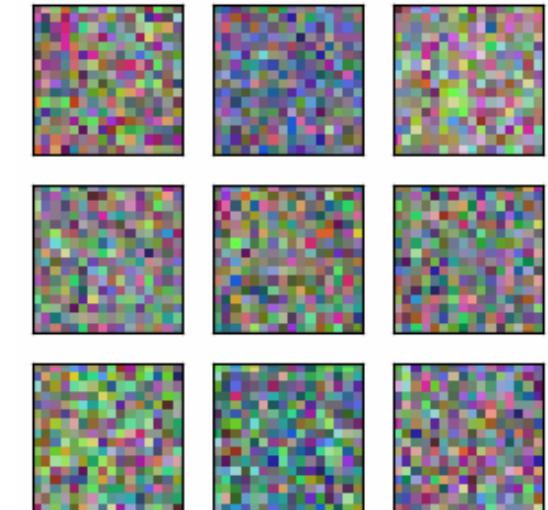
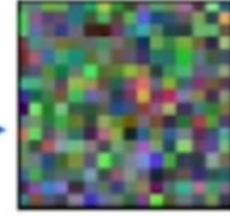
Diffusion Inference

Noise sample



Trained NN

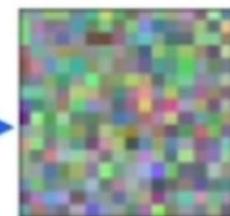
Predicted noise



Generation



Trained NN

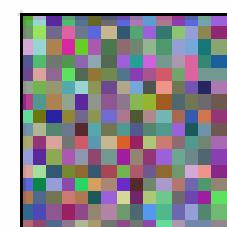
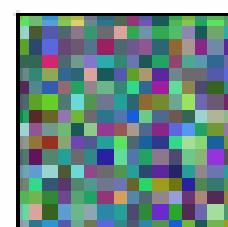
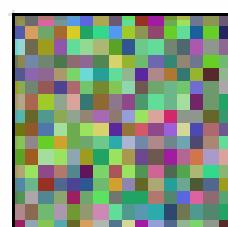
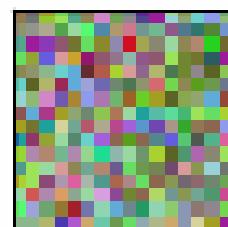
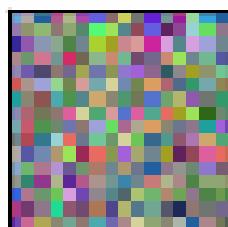
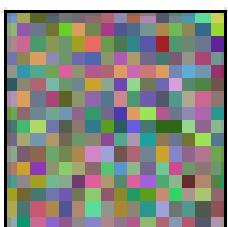
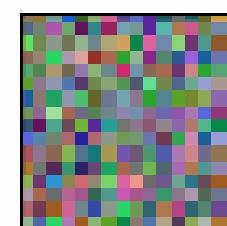
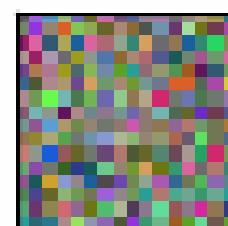
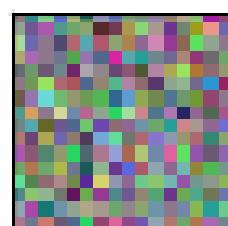
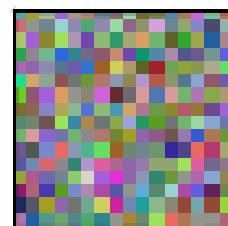
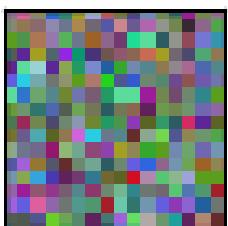
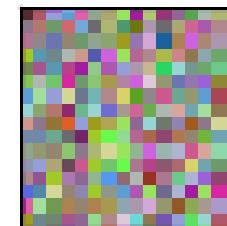
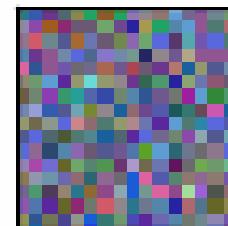
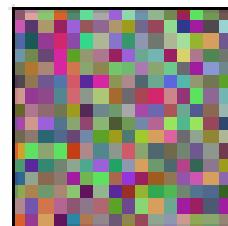
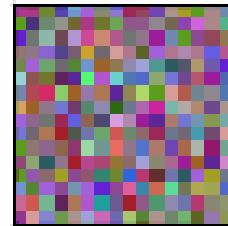
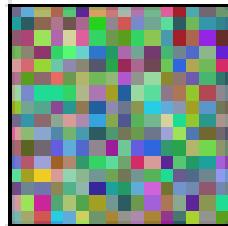
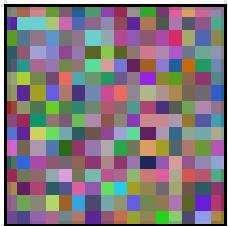


⋮

Diffusion Inference

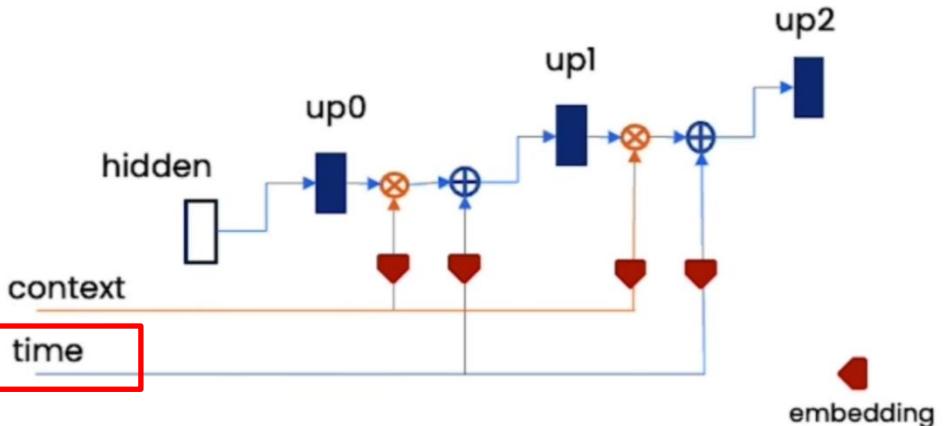
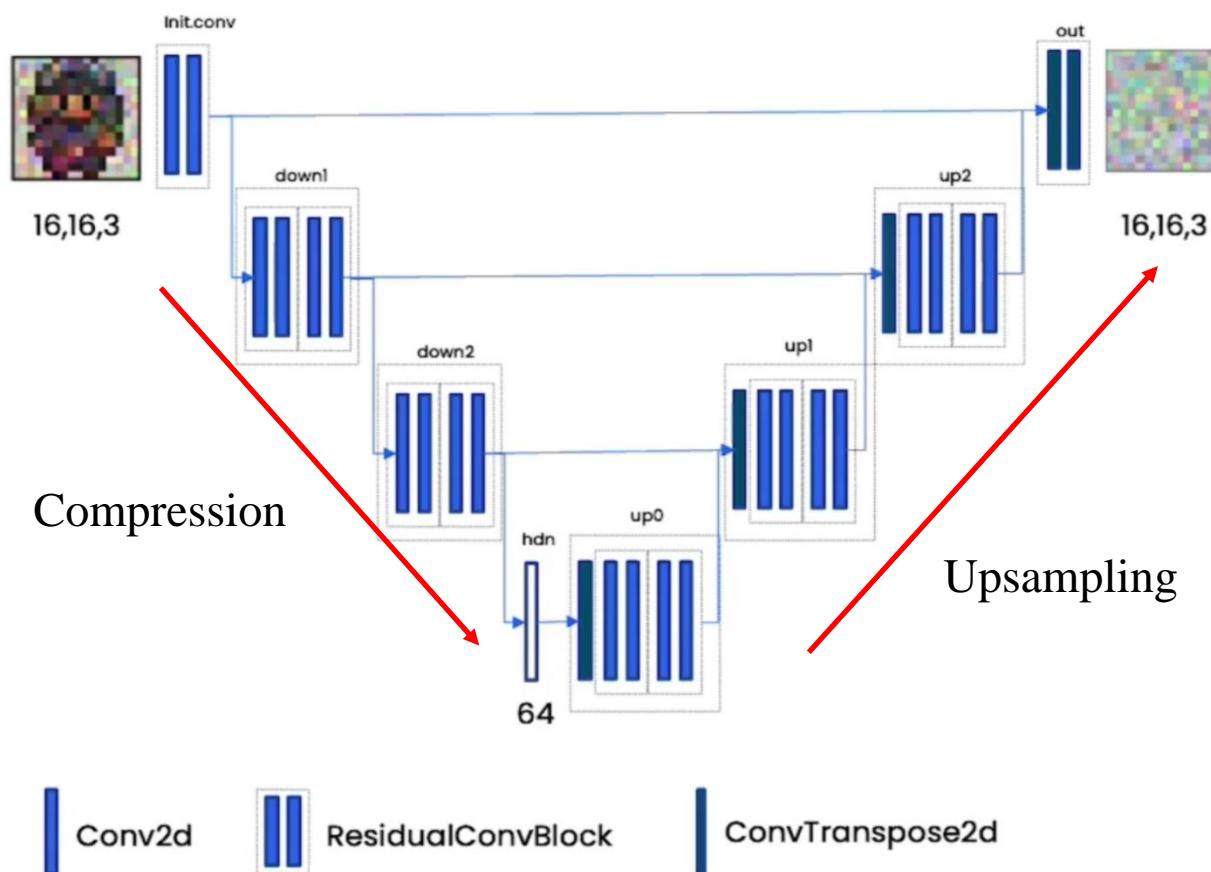
- Trick/Empirically

- Add in additional noise before it gets passed to the next step
- Will not collapse to something closer to the average of the dataset



Neural Network

- Input & outputs have the same shape.
- Can take additional inputs.



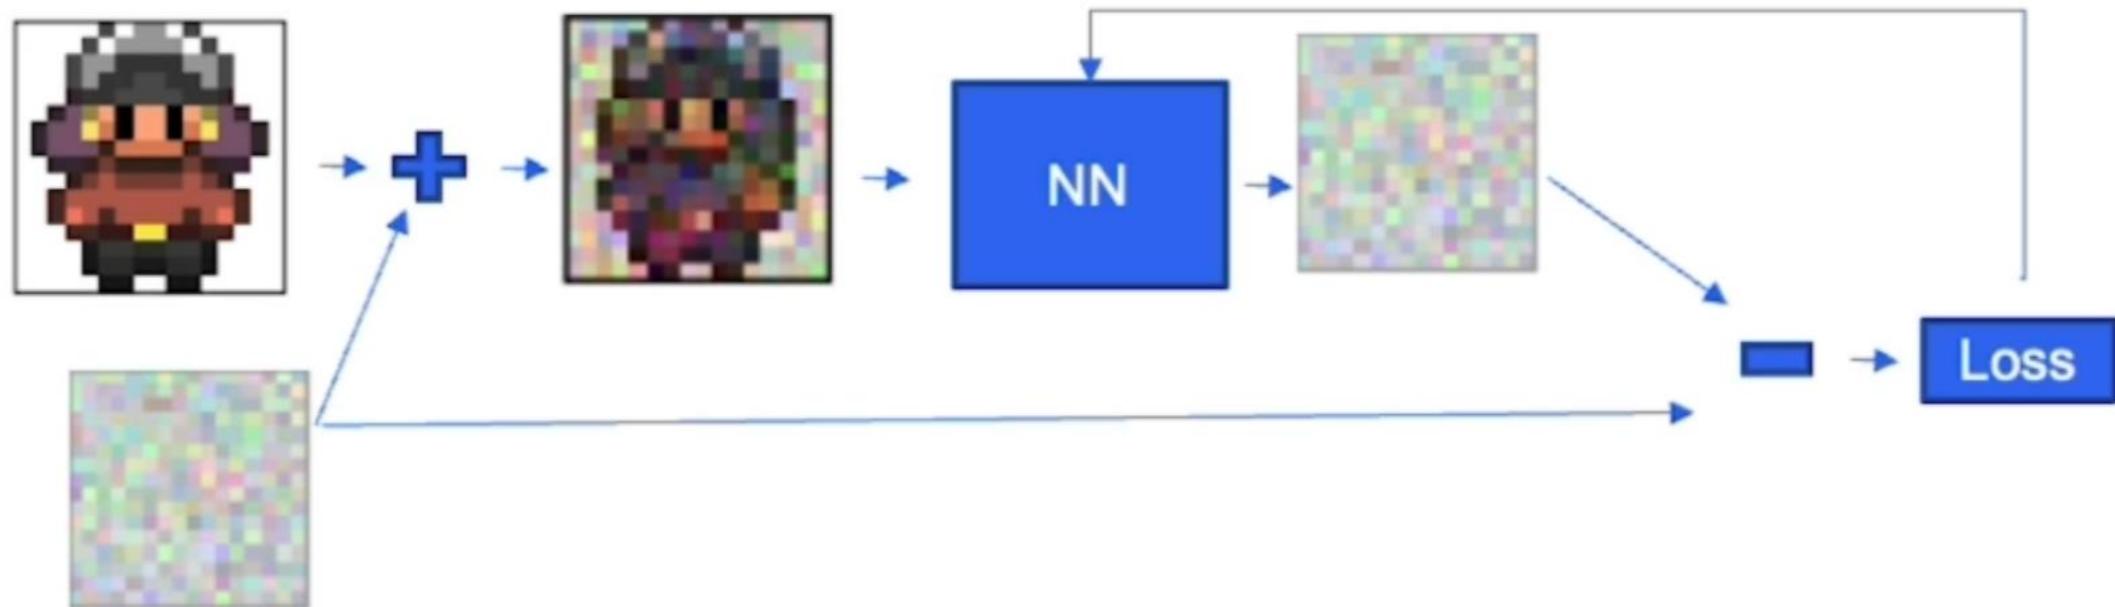
```

# embed context and timestep
cemb1 = self.contextembed1(c).view(-1, self.n_feat * 2, 1, 1)
# (batch, 2*n_feat, 1,1)
temb1 = self.timeembed1(t).view(-1, self.n_feat * 2, 1, 1)
cemb2 = self.contextembed2(c).view(-1, self.n_feat, 1, 1)
temb2 = self.timeembed2(t).view(-1, self.n_feat, 1, 1)

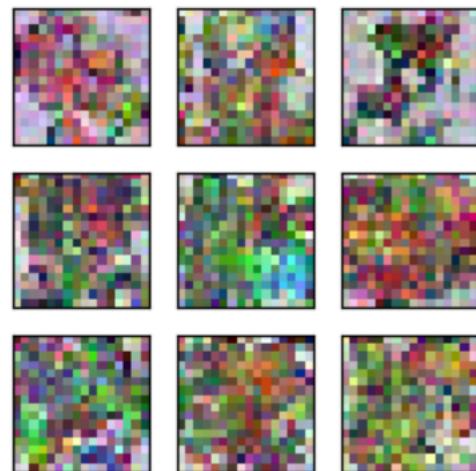
up1 = self.up0(hiddenvec)
# add and multiply embeddings
up2 = self.up1(cemb1*up1 + temb1, down2)
up3 = self.up2(cemb2*up2 + temb2, down1)
out = self.out(torch.cat((up3, x), 1))
return out
  
```

Training

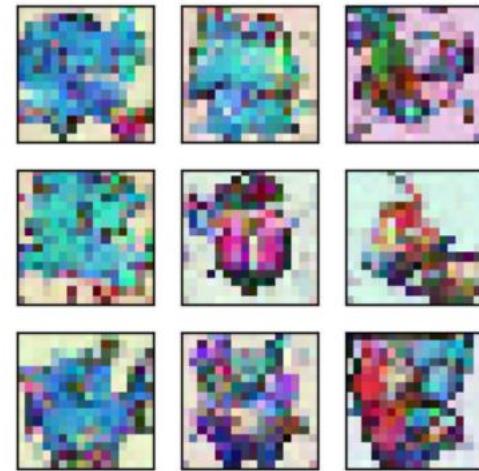
- Neural network learns to predict noise
 - The distribution of what is not noise
 - Sample random timestep (a.k.a noise level) per image to train more stably



Training



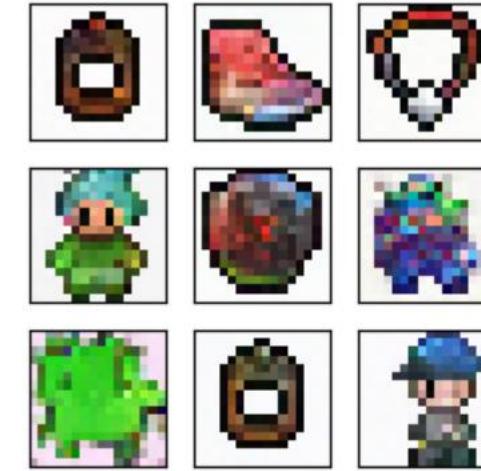
Epoch 0



Epoch 4



Epoch 8



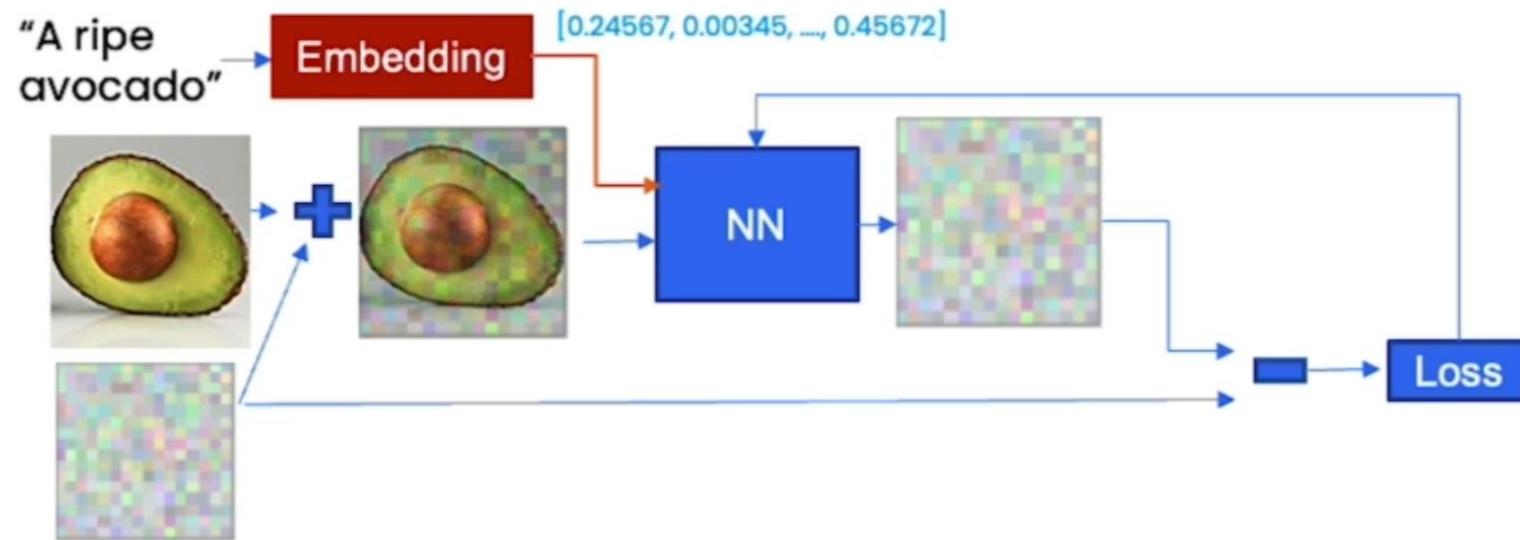
Epoch 31

Controlling



Embedding vector captures meaning

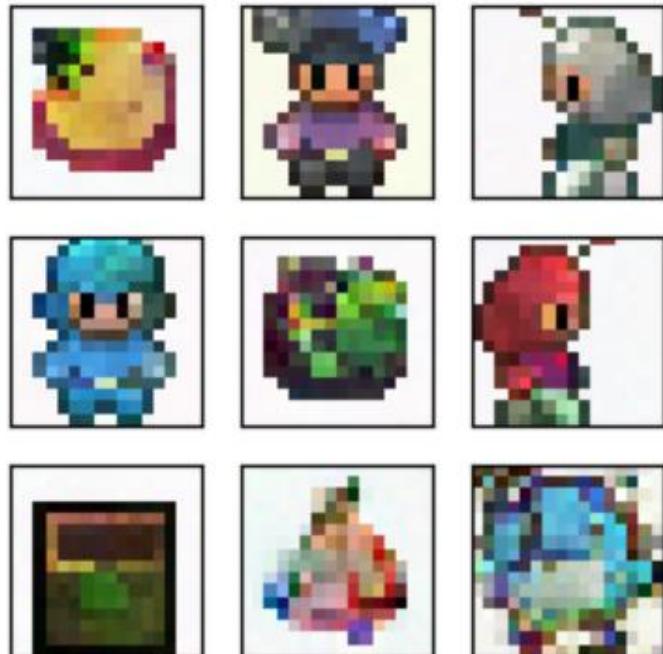
Training



Controlling

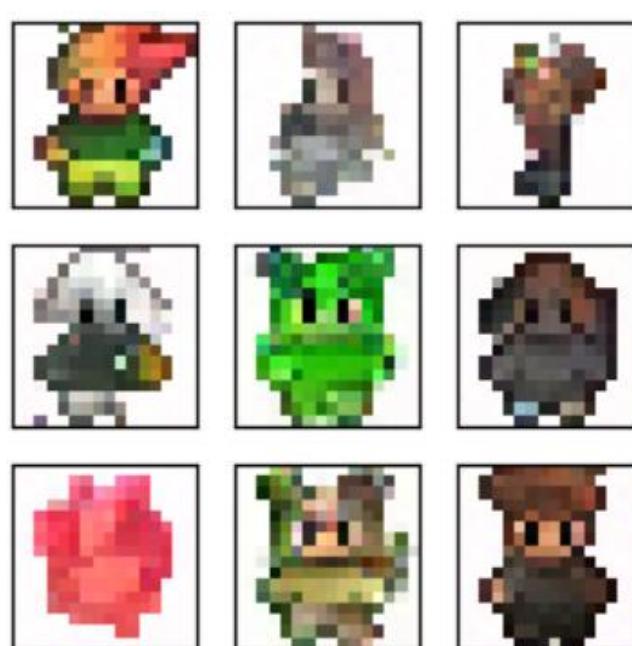
- Context Definition:

- hero, non-hero, food, spell, side-facing



Input Context:

Food	Hero	Side-facing
Hero	Food	Side-facing
Spell	Food	No-hero



Combination:

Food	Hero	Side-facing
Hero	Food	Side-facing
Spell	Food	No-hero

```
[1,0,0,0,0],  

[1,0,0.6,0,0],  

[0,0,0.6,0.4,0],  

[1,0,0,0,1],  

[1,1,0,0,0],  

[1,0,0,1,0],  

[1,0,1,0,0],  

[1,1,0,0,0],  

[1,0,0,1,0],
```

Homework

Diffusion Demo

- **Project #1 (For everyone)**

- **Introduction:** explore the training, sampling and controlling in the diffusion models.
- **Implementation:**
[How-Diffusion-Models-Work](#)
- **Tasks:**
 - ❑ Training your diffusion models.
 - ❑ Explore the hyper-parameter in the diffusion models.
 - ❑ Sampling algorithm: DDPM, DDIM.
 - ❑ Control strategy.
 - ❑
- **Members:** for everyone.



Contents

- Introduce to Pytorch
- Project Guidelines
 - Project Expectations
 - Diffusion Demo
 - Picking a project idea
 - Final Report

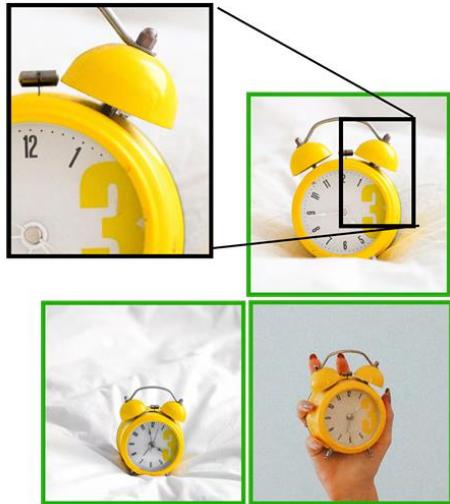
Picking a project idea

- Do what is interesting to you.
- Do not directly what seems easiest.
- Considerations:
 - Data:
 - ✓ Existing data for your problem.
 - ✓ Need to collect and preprocess data by yourself.
 - Code:
 - ✓ Based on an existing implementation.
 - ✓ Implement from scratch.

Picking a project idea

- **Attachable Project (Dreambooth)**

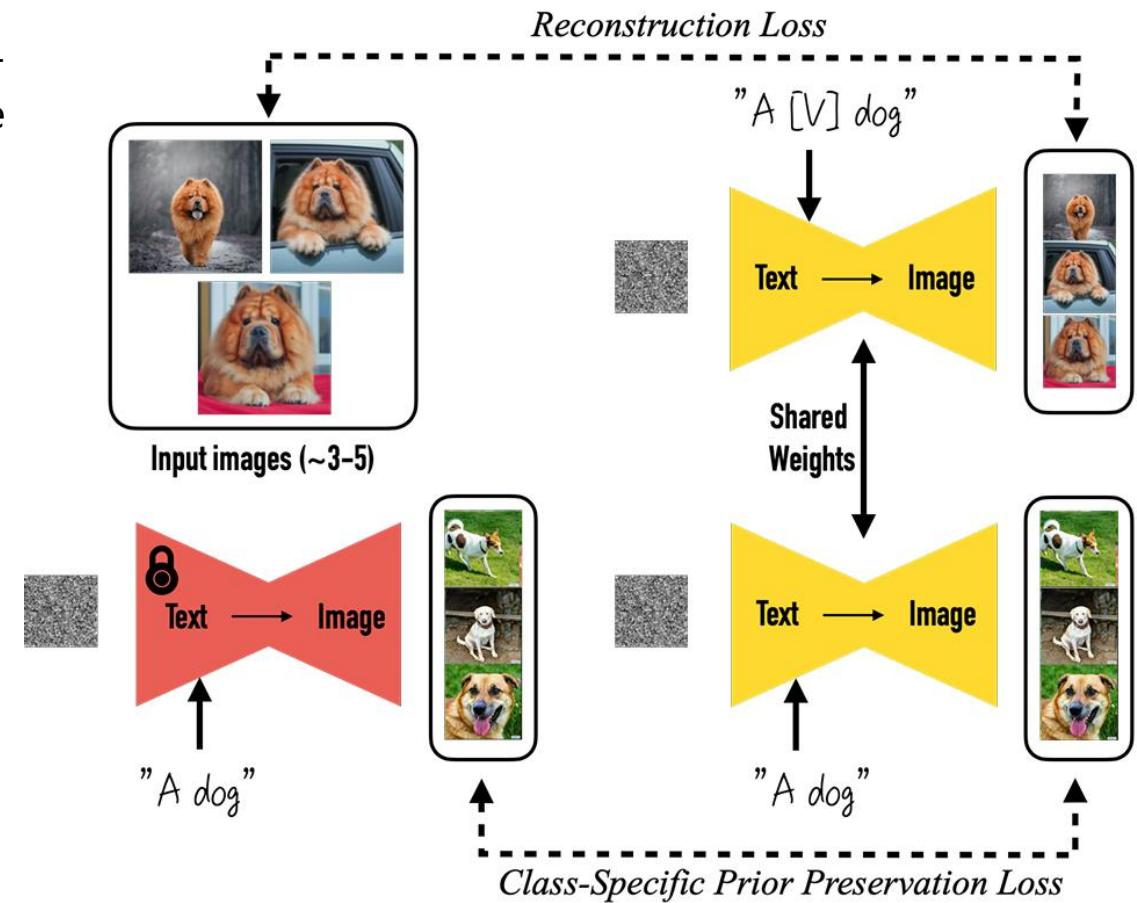
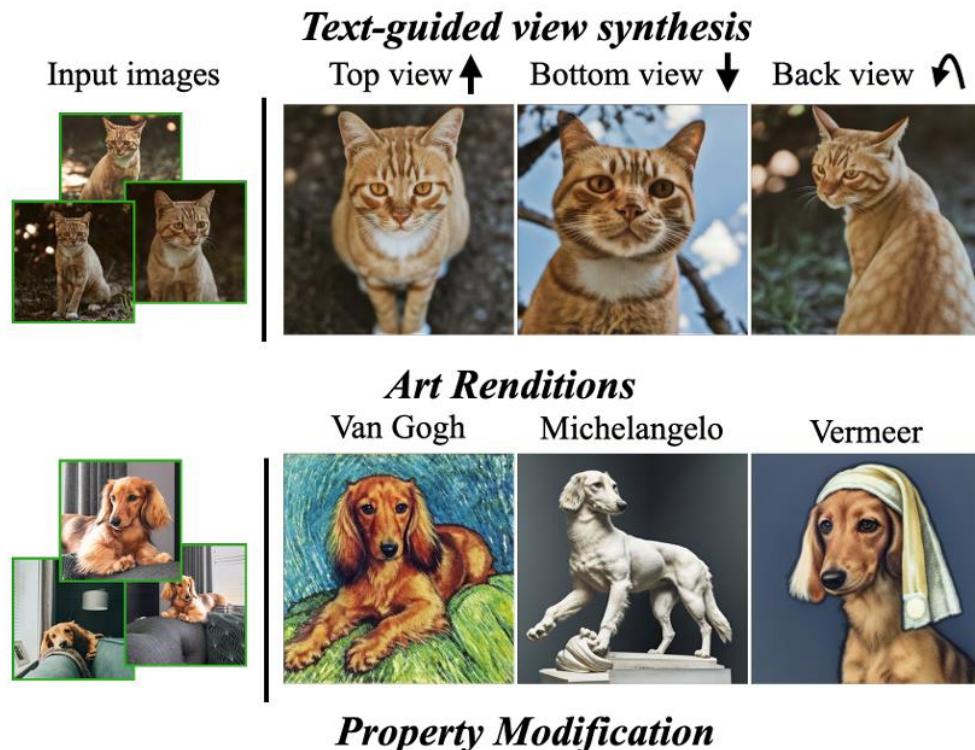
- Text Prompts (DALL-E2/Imagen): retro style yellow alarm clock with a white clock face and a yellow number three on the right part of the clock face in the jungle
- Dreambooth: a [V] clock in the jungle



Picking a project idea

- Attachable Project (Dreambooth)

- **Introduction:** Given 3-5 images, fine-tuning a text-to-image diffusion model with a unique identifier [V] and the name of class.



Dreambooth (Ruiz, ICCV2023)

Picking a project idea

- Attachable Project (Dreambooth)

Color modification (“*A [color] [V] car*”)



Input



purple

red

yellow

blue

pink

Hybrids (“*A cross of a [V] dog and a [target species]*”)



Input



bear

panda

koala

lion

hippo

Picking a project idea

- **Attachable Project (Dreambooth)**

- **Introduction:** Given 3-5 images, fine-tuning a text-to-image diffusion model with a unique identifier [V] and the name of class.

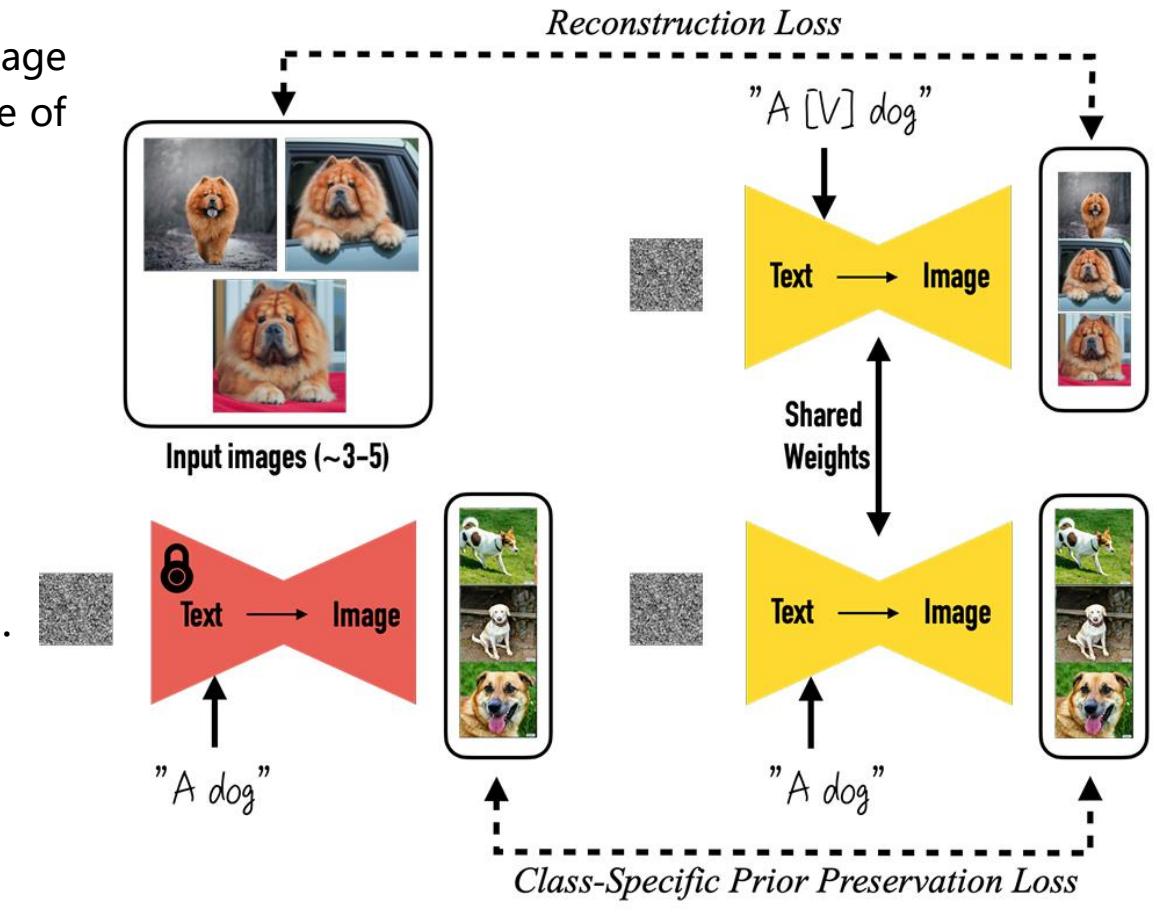
- **Implementation:**

[google/dreambooth \(github.com\)](https://github.com/dreambooth)

- **Tasks:**

- Build your own fine-tuning datasets.
- Explore the training process of diffusion model.
- Analysis the inferences of your models.
- How to select the identifier [V].
-

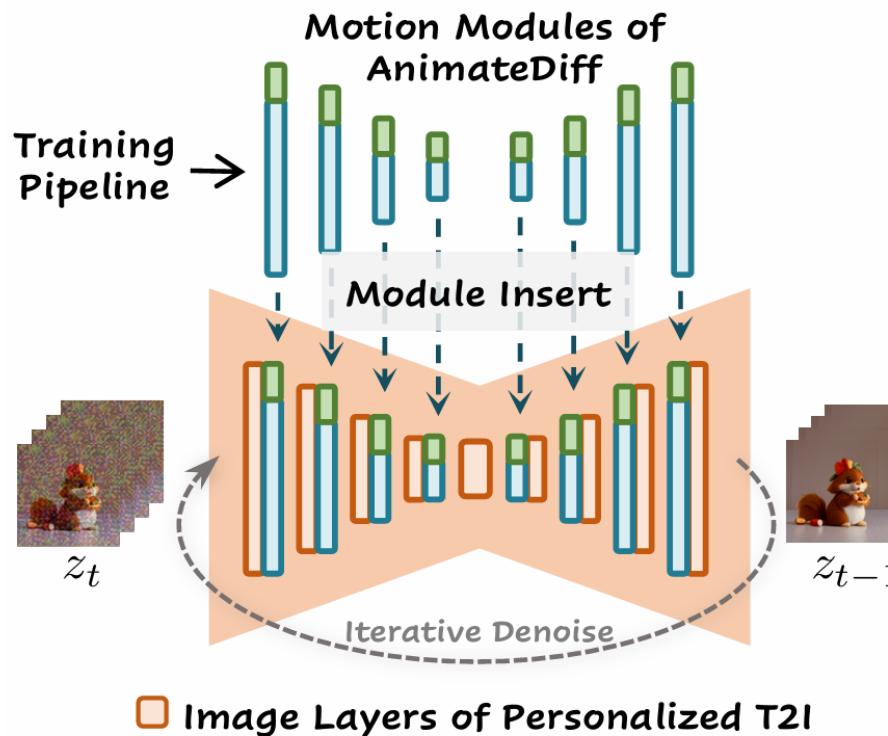
- **Members:** no more than 3.



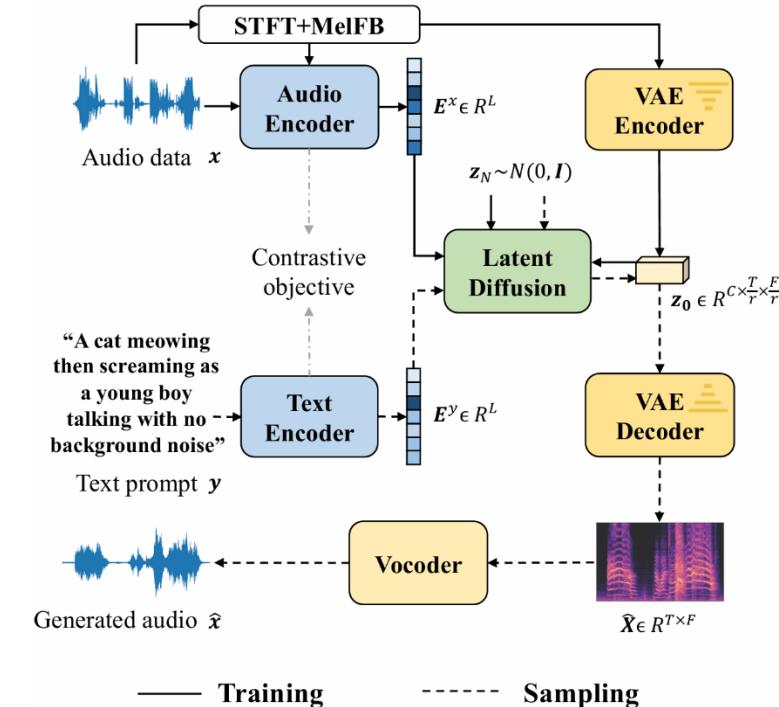
Dreambooth (Ruiz, ICCV2023)

Picking a project idea

- Project based on your interests (no more than 6 members)

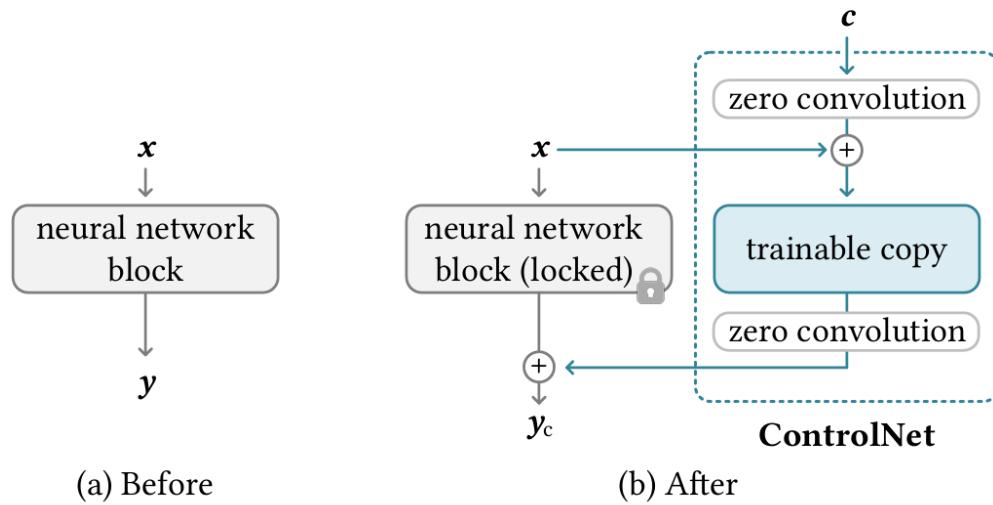


Video Generation
[AnimateDiff \(Guo, ICLR2024\)](#)

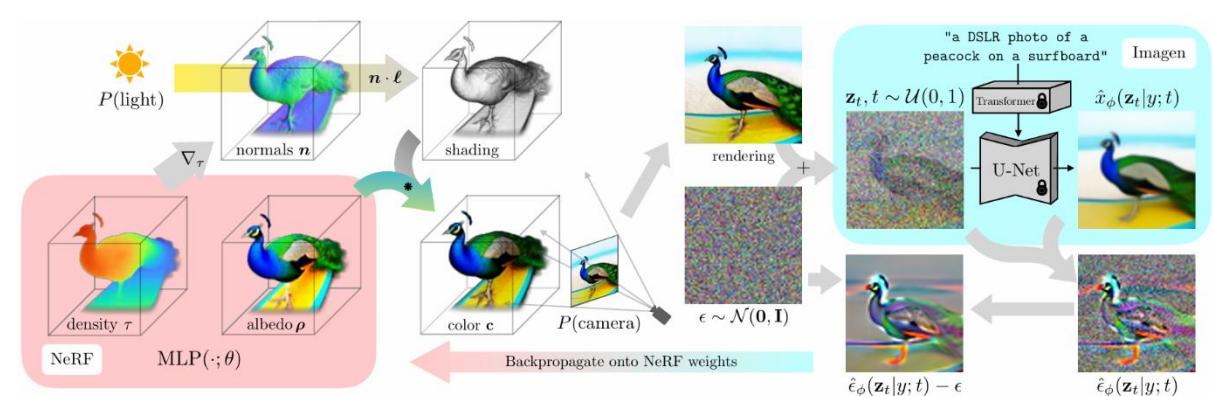


Picking a project idea

- Project based on your interests (no more than 3 members)



Controllable Diffusion
[ControlNet \(Zhang ICCV 2023\)](#)



Picking a project idea

- Literature review about diffusion models (algorithm or applications)
 - **Contents**
 - Introduction:
 - Preliminary
 - Literature Review (applications in specific domains)
 - Key Challenges
 - Future Direction
 -
 - Conclusions
 - **Members:** only for 1.

Contents

- Introduce to Pytorch
- Project Guidelines
 - Project Expectations
 - Diffusion Demo
 - Picking a project idea
 - Final Report

Course Schedule

Lecture	Topic	Content
1	Introduction	Overview of this Course
2	Basics of Generative AI (1)	VAE, GAN, Flow-based models, and applications
3	Basics of Generative AI (2)	Probability distributions, random variables
4	Fundamentals of Diffusion Models (1)	Diffusion model principles, DDPMs, DDIMs, SGMs, Score SDEs, VDMs
5	Fundamentals of Diffusion Models (2)	Model structures for generation processes: Unet, DiT, Latent space
6	Machine Learning Frameworks and Tools	Machine learning frameworks and tools: training, inference, optimization
7	Images + Diffusion Models	LDM、Stable Diffusion、DALL-E、GigaGAN
8	Audio + Diffusion Models	Audio Diffusion Models、VideoDiffusion、Sora
9	3D + Diffusion Models	NeRF、3D-VAE、DreamFusion
10	Biology (Structure) + Diffusion Models	AlphaFold3、ESMFold、RFdiffusion、SE(3) Diffusion
11	Physics and Meteorology + Diffusion Models	GraphCast、Pangu-Meteorology、NowcastNet、Fuxi
12	Advanced Topics in Diffusion Models (1)	External expert talk
13	Advanced Topics in Diffusion Models (2)	External expert talk
14	Paper and Project Presentations	
15	Paper and Project Presentations	
16	Paper and Project Presentations	

Final Report

- **Report Proposal**

- **Brief Introduction**
- **Problems**
- **Methods**
- **Members**
- **Framework** (optional in this stage)

- **Report Proposal**

- **Slides**
- **Reports**
- **Posters**

- **More details in the course homepage**

Lecture	Topic
1	Introduction
2	Basics of Generative AI (1)
3	Basics of Generative AI (2)
4	Fundamentals of Diffusion Models (1)
5	Fundamentals of Diffusion Models (2)
6	Machine Learning Frameworks and Tools
7	Images + Diffusion Models
8	Audio + Diffusion Models
9	3D + Diffusion Models
10	Biology (Structure) + Diffusion Models
11	Physics and Meteorology + Diffusion Models
12	Advanced Topics in Diffusion Models (1)
13	Advanced Topics in Diffusion Models (2)
14	Paper and Project Presentations
15	Paper and Project Presentations
16	Paper and Project Presentations



Thanks!